

Stereo Reconstruction

Alex Rasla

UC Santa Barbara

alexrasla@ucsb.edu

1 Algorithms

In order to achieve a good stereo reconstruction, I first ensured I could find an appropriate window size, compute the best raw sum of squared differences (SSD), and then generate a baseline disparity map. Once I tested this baseline disparity generation on a variety of images, I decided to implement the Maximum Likelihood Stereo Algorithm [1] to further enhance the quality of my disparity maps. This algorithm takes a Disparity Space Image (DSI) of a scanline between two stereo images and uses dynamic programming to get the shortest path (lowest sum of dissimilarity scores) from the first to the last column. Once this shortest path is calculated, we can choose to include occlusion filling to further increase the quality of the disparity maps.

2 Data

For this stereo reconstruction project, I used the 2006 Middlebury Stereo Dataset [2]. This dataset is a very trusted and standard dataset for stereo reconstruction algorithms because of its high-quality stereo images and accurate ground truth disparity maps. Specifically, I utilized the Aloe, Cloth1, Cloth2, Midd1, Midd2, Rocks1, and Wood1 datasets for both the full-size (1282x1110) and half-size (665x555) images. Unfortunately, because it is very difficult and resource dependent to generate accurate ground truth disparity maps, I was unable to create my own dataset to test my stereo reconstruction algorithm.

3 Libraries

For this assignment, I used OpenCV, NumPy, and Numba. I used OpenCV to simply read and write images into memory, since my stereo matching algorithm did not require any feature matching or other low-level optimizations. Rather, it required a lot of pixel processing, computing for-

mulas between windows, and matrix operations, all of which I used NumPy for. Furthermore, to speed up NumPy computation and vectorize the code, I used Numba, a high performance Python compiler that converts NumPy loops and dynamic programming into very fast machine code.

4 Implementation Details

In order to run this program, simply type the command:

```
python3 stereo.py --dir [head directory of images]
                  --left [left image]
                  --right [right image]
                  --out [output directory]
                  --scale [disparity scale]
```

For this command, the `--dir` flag should be the head directory of images that contains subdirectories with the downloaded/test datasets from Middlebury [2]. The program goes through all the subdirectories in `--dir`, gets right and left disparity images, saves left and right disparity maps into the out folder in the subdirectory, and generates a bar graph of the Bad Matched Pixels (BMP) [3] and the runtime to calculate the maps for each of the different datasets. For the full-size images (1282x1110), the scale should be 1; for half-size, the scale should be 2; for third-size, it should be 3.

5 Results

The results of my stereo reconstruction produced visually consistent disparity maps for all datasets, but produced expected results based on the quantitative BMP evaluation metric.

$$BMP = \frac{1}{N} \sum_{(x,y)} (|d_C(x,y) - d_T(x,y)| > \delta_d)$$

The BMP metric is a binary metric that checks whether a resultant disparity map is within a given

	$\delta = 1$	$\delta = 5$	$\delta = 20$
Full Size	0.41-0.77	0.40-0.76	0.40-0.75
Half Size	0.53-0.69	0.52-0.67	0.52-0.66

Table 1: The range of left image stereo reconstruction Bad Matched Pixel values for Dyanmic Programming with Occlusion Filling




Figure 1: **1a** is the raw SSD left disparity map for the Cloth1 full-size dataset; **1b** is the result of dynamic programming without occlusion filling; **1c** is dynamic programming with occlusion filling; **1d** is the ground truth; **1e** is the original RGB image

δ value, and then computes the percentage of those that are within the value. In the original paper [3], the authors chose a δ value of 1, meaning they only allowed a difference of 1 pixel between the generated disparity map and the ground truth. In my evaluation, I focused on δ values of 5.

My initial, baseline attempt at generating disparity maps was simply using SSD. This disparity map was structurally accurate, but evaluated to a BMP range of 0.55 – 0.76 (full-size, $\delta = 5$) for each datasets left disparity maps. Next, to improve the BMP values and quality of the disparity maps, I used dyanmic programming without occlusion filling which significantly narrowed the BMP range to between 0.55 – 0.68 for the left disparity maps (full-size, $\delta = 5$). Finally, when I added occlusion filling, this again improved the BMP value to between 0.40 – 0.76 (full-size, $\delta = 5$) and 0.52 – 0.67 (half-size, $\delta = 5$). Table 1 shows the range of BMP values for dynamic programming with occlusion filling, Figure 1 shows a visual comparison of the different algorithms used to generate disparity maps, and Figure 2 and Figure 3 show bar graph plots of the the different BMP values and the runtime to generate them.




Figure 2: A plot of the BMP score for the left and right disparity maps for **full-size** images in the Aloe, Cloth1, Cloth2, Midd1, Midd2, Rocks1, and Wood1 datasets

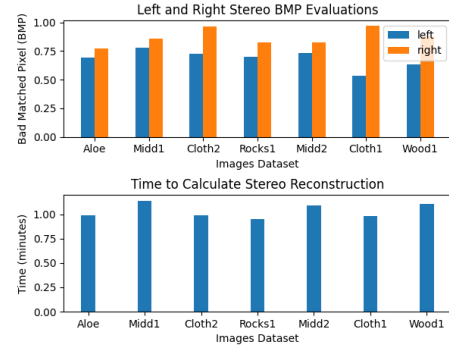


Figure 3: A plot of the BMP score for the left and right disparity maps for **half-size** images in the Aloe, Cloth1, Cloth2, Midd1, Midd2, Rocks1, and Wood1 datasets

6 Discussion

Throughout this project, I explored the difference between SSD disparity maps, dynamic programming disparity maps, and dynamic programming with occlusion filling disparity maps. Because the latter two methods are built upon the standard SSD scanline matching, I had to ensure the SSD was computed efficiently and quickly. To do this, I first used NumPy array slicing to create a Disparity Space Image (DSI) between the left scanline and right scanline. This DSI gave me a $M \times M$ matrix, where each value represented the SSD between a left pixel (row) and right pixel (column) on the scanline. Once this DSI was calculated, to generate the SSD disparity map, I simply took the maximum values along each row to get the best matching (lowest SSD) right pixel for every left pixel.

To improve upon this and achieve a better result, I implemented the dynamic programming algorithm described in [1]. In order to implement this algorithm, it required an occlusion cost parameter which, after experimentation, I found generated the best BMP values at 100,000. This occlusion cost represents the cost to “occlude” a DSI matrix value, and proceed to the next pixel value. This algorithm enhanced qualitative visual accuracy of the raw SSD results by finding the shortest path on a particular scanline and omitting SSD values that were greater than the occlusion cost. Typically, I noticed that with an occlusion cost of any multiple of 100,000, the BMP values increased. This makes sense because the shortest path algorithm relies on the occlusion cost to omit outlier pixels that have SSD window values that are too high. If these values aren’t omitted, poor disparity maps will be generated. On the other hand, if the cost is too low, we may end up with the same result as the SSD image, making the optimizations irrelevant. To further improve upon the dynamic programming algorithm, I implemented occlusion filling, which simply uses the disparity value of an adjacent pixel as opposed to giving it a grayscale value of 0.

Alongside experimenting with the occlusion cost and reconstruction algorithms, I also experimented with the window size and the magnitude of δ . When changing the window size from 10 to 30, the runtime was around 2.5 times slower and the BMP range of values increased significantly to between 0.52 – 0.89 (half-size, $\delta = 5$). For Cloth1 specifically, the BMP increased from 0.52 to 0.74, a

nearly 25% increase poorly matched pixels. This is because larger window sizes cannot accurately detect details in an image, causing a pixel by pixel evaluation metric like BMP to be heavily impacted. When increasing the delta values even slightly, I expected the BMP values to decrease and lead to a better quantitative evaluation of the disparity images; however, as we can see in Table 1, the change in δ values did not produce a significant change in BMP. This is likely because the values that lie outside of δ lie extremely far outside of δ and require δ values of at least a two orders of magnitude greater to effect the BMP score.

References

- [1] Ingemar J. Cox et al. “A Maximum Likelihood Stereo Algorithm”. In: *Computer Vision and Image Understanding* 63.3 (1996), pp. 542–567. ISSN: 1077-3142. DOI: <https://doi.org/10.1006/cviu.1996.0040>. URL: <https://www.sciencedirect.com/science/article/pii/S1077314296900405>.
- [2] Heiko Hirschmüller and Daniel Scharstein. “Evaluation of Cost Functions for Stereo Matching”. In: *2007 IEEE Conference on Computer Vision and Pattern Recognition* (2007), pp. 1–8.
- [3] Daniel Scharstein and Richard Szeliski. “A Taxonomy and Evaluation of Dense Two-Frame Stereo Correspondence Algorithms”. In: *Int. J. Comput. Vision* 47.1–3 (Apr. 2002), pp. 7–42. ISSN: 0920-5691. DOI: [10.1023/A:1014573219977](https://doi.org/10.1023/A:1014573219977). URL: <https://doi.org/10.1023/A:1014573219977>.